

Jimmy Huang

📍 Taipei, Taiwan
 ✉️ skyrocketqy81@gmail.com
 🌐 rendercv.com
 🔗 rendercv
 🐙 rendercv

Welcome to RenderCV

RenderCV reads a CV written in a YAML file, and generates a PDF with professional typography.

Each section title is arbitrary.

You can choose any of the 9 entry types for each section.

Markdown syntax is supported everywhere. This is **bold**, *italic*, and [link](#).

Education

MS	Chung Yuan Christian University , Applied Mathematics	Sept 2020 – June 2022
BS	Chang Jung Christian University , Accounting and Information	Sept 2015 – June 2019

Experience

Microsoft Research , Research Intern	Redmond, WA
- Implemented novel self-supervised learning framework for low-resource language modeling - Research integrated into Azure Cognitive Services, reducing training data requirements by 60%	May 2019 – Aug 2019 4 months

Projects

FlashInfer	Jan 2023 – present
Open-source library for high-performance LLM inference kernels	
- Achieved 2.8x speedup over baseline attention implementations on A100 GPUs - Adopted by 3 major AI labs, 8,500+ GitHub stars, 200+ contributors	

Selected Honors

- MIT Technology Review 35 Under 35 Innovators (2024)
- Forbes 30 Under 30 in Enterprise Technology (2024)
- ACM Doctoral Dissertation Award Honorable Mention (2023)
- Google PhD Fellowship in Machine Learning (2020 – 2023)
- Fulbright Scholarship for Graduate Studies (2018)

Skills

Languages: Python, C++, CUDA, Rust, Julia

ML Frameworks: PyTorch, JAX, TensorFlow, Triton, ONNX

Infrastructure: Kubernetes, Ray, distributed training, AWS, GCP

Research Areas: Neural architecture search, model compression, efficient inference, multi-agent RL

Patents

- Adaptive Quantization for Neural Network Inference on Edge Devices (US Patent 11,234,567)
- Dynamic Sparsity Patterns for Efficient Transformer Attention (US Patent 11,345,678)
- Hardware-Aware Neural Architecture Search Method (US Patent 11,456,789)

Invited Talks

4. Scaling Laws for Efficient Inference — Stanford HAI Symposium (2024)
3. Building AI Infrastructure for the Next Decade — TechCrunch Disrupt (2024)
2. From Research to Production: Lessons in ML Systems — NeurIPS Workshop (2023)
1. Efficient Deep Learning: A Practitioner's Perspective — Google Tech Talk (2022)